

Analysis and evaluation of audio-similarity algorithms for cover and live song identification



Martin Angelov

MSci Computer Science with Industrial Year, Student ID 1422087

School of Computer Science, University of Birmingham

Supervisor

Prof. Achim Jung

Acknowledgements

Write here your acknowledgements...

Abstract

The abstract goes here. The abstract should be self-contained and:

- clearly state the problem dealt with by the thesis;
- give a synthetic description of the proposed solution;
- highlight the sense in which the proposed solution enhances the state of the art.

Contents

1	Introduction	1
1.1	Project overview	1
1.2	Report structure	3
2	Background theory	4
2.1	Basic properties of audio signal	4
2.1.1	Test subsection	5
2.2	Audio transformation techniques	5
3	Related work	6
3.1	Examination of other audio similarity techniques and algorithms not analysed by the project	6
3.2	Scientific paper rewritten	6
4	The task	7
4.1	Design	7
4.2	Evaluation methods and metrics	7
4.3	Datasets used for evaluation	7
5	The algorithms	8
5.1	Osmalskyj big algorithm	9

5.2	Osmalskyj weak features	9
5.3	Ellis cross-correlation algorithm	9
5.4	Osmalskyj quantisation algorithm	9
5.5	Tralie timbre algorithm	9
5.6	Rank aggregation techniques	9
5.7	Rafii audio fingerprinting algorithm	9
6	The benchmark	10
6.1	Implementation details	10
6.2	Brief usage information	10
6.3	Algorithm structure in the benchmark	10
6.4	Result format produced by benchmark	10
7	Results	11
7.1	Best results	11
7.2	Comparison to results from papers	11
7.3	Result analysis	11
8	Further work	12
9	Challenges	13
9.1	Lack of datasets	13
9.2	Lack of universal comparison metric?	13
9.3	Academic papers algorithm description	13
10	Project management	14
10.1	Using GitLab	14
10.2	Canvas logs	14
10.3	other? Gantt chart?	14

CONTENTS

11 Conclusion	15
References	17

List of Figures

List of Tables

Chapter 1

Introduction

1.1 Project overview

Music information retrieval (MIR) is an area of analysis dedicated to extracting information from music. It combines many different disciplines of science including psychology, psychoacoustics, signal processing and computer science. One of the main aims when applying MIR techniques solving the task of song identification, i.e. matching an audio stream to a particular song [1]. This is usually achieved through a form of hashing applied on the digital signal and comparing the resulting representation to a reference fingerprint [2], [3]. This approach returns good results for the task, since we can easily quantify a good match between both fingerprints.

We can further modify the original song identification task to apply to cover songs. A cover song is a very creative reinterpretation of a released song usually performed by an artist different than the original. The cover can therefore differ significantly from the origin in tempo, pitch or song structure (*add more*). The amount of variation in a cover strongly depends on the genre of the primary track - Western popular music pieces are for example more likely to be transformed

than ones from classical music [4]. Therefore the only remaining common feature between the cover song and the original is the underlying fundamental melody of the piece and potentially the lyrics.

Because of these potential disparities between two versions of a single song, the problem of identifying covers of songs is much more difficult than determining an identical match with the original. The above fingerprinting approach has been attempted [5] and the results are insignificant [6]. Direct comparison between the fingerprints of the song is unable to capture the remaining similarity within two audio files. Other MIR methods need to be considered in order to measure similarity when attempting cover song recognition.

The general advances of technology have allowed companies such as Spotify [7], Apple [8], SoundCloud [9] and more to create large-scale music databases and offer them as commercial services. Proportionally to the increasing availability of large music collections grows the need for managing the volumes of audio information through MIR techniques, with cover song identification being one of them. As a consequence most modern mechanisms to cover song recognition work by comparing an audio track called *query song* against a large database of songs, a *reference database*. Each mechanism is evaluated based on its similarity estimation performance, as well as its scalability as we increase the database size.

This project analyses the principles of a set of non-hashing based cover song identification algorithms and evaluates their performance. Most of the examined algorithms are designed to work with large-scale databases and follow the workflow model described above. The evaluation considers only their similarity estimation results and does not account for scalability. After analysis of the results a hypothesis on the best performing audio similarity technique is established (*or maybe devised?*).

1.2 Report structure

The sections of this report are as follows:

- *Chapter 2* offers a summary of the background information required to understand and implement the audio similarity algorithms
- *Chapter 3* explores other state of the art methods of measuring similarity not examined in detail by the project
- *Chapter 4* provides a description of the evaluation task through which each algorithm is analysed
- *Chapter 5* contains detailed descriptions of each algorithm
- *Chapter 6* expands on implementation details related to the benchmark tool
- *Chapter 7* outlines the best results achieved and offers an analysis on them
- *Chapter 8* summarises potential further contributions to the project
- *Chapter 9* discusses the main challenges related to the project and the task of cover song identification
- *Chapter 10* is a summary of the project management techniques utilised during the project

Chapter 2

Background theory

Each type of information extracted from an audio stream is referred to as an *audio feature*. Audio features are mainly derived using various transformations on the signal based on some basic properties of sound. This section presents low-level theory required to understand the high-level description of how each feature is obtained further in the report.

2.1 Basic properties of audio signal

A *digital audio signal* is a representation of the continuous sound wave as a discrete series of binary numbers. This representation helps preserve the *frequency* (the speed of the vibrations), as well as the *amplitude* (the fluctuations of the vibrations) of the sound. The energy that each sound wave emits through vibrations is called *sound energy* and its rate is measured through *sound power* [10]. The majority of audio features use frequency or power as a primary audio property used to define the feature (*modify/change*).

The process of converting an analog sound wave to a digital one involves a process of extracting points (samples) from the continuous signal and using them

2.2 Audio transformation techniques

to describe the signal into a discrete form. This method is called *sampling* and the amount of samples collected per time frame is *sample rate*. The representation of a song used during feature extraction is a sequence of samples extracted from the digital signal of the song based on its sample rate.

2.1.1 Test subsection

2.2 Audio transformation techniques

Chapter 3

Related work

Summary

short summary of the chapter...

One or more chapters should be devoted to the description of the proposed approach...

In particular, this chapter describes the design adopted by this research to achieve the aims and objectives stated in the Introduction.

3.1 Examination of other audio similarity techniques and algorithms not analysed by the project

3.2 Scientific paper rewritten

Chapter 4

The task

Summary

Discuss here the methodology used in the study, the stages by which the methodology was implemented, and the research design; For examples, one section details the participants in the study, another section lists all the instruments used in the study and justifies their use; another section outlines the procedure (algorithms, code,..) used; a section discusses how the data was analysed, etc..

4.1 Design

4.2 Evaluation methods and metrics

4.3 Datasets used for evaluation

Chapter 5

The algorithms

Summary

Details all the results of your study here (exploits graphics for results visualization). This chapter should also contain a full discussion, interpretation and evaluation of the results.

- 5.1 Osmalskyj big algorithm
- 5.2 Osmalskyj weak features
- 5.3 Ellis cross-correlation algorithm
- 5.4 Osmalskyj quantisation algorithm
- 5.5 Tralie timbre algorithm
- 5.6 Rank aggregation techniques
- 5.7 Rafi audio fingerprinting algorithm

Chapter 6

The benchmark

Summary

Details all the results of your study here (exploits graphics for results visualisation). This chapter should also contain a full discussion, interpretation and evaluation of the results.

6.1 Implementation details

6.2 Brief usage information

6.3 Algorithm structure in the benchmark

6.4 Result format produced by benchmark

Chapter 7

Results

Summary

Details all the results of your study here (exploits graphics for results visualisation). This chapter should also contain a full discussion, interpretation and evaluation of the results.

7.1 Best results

7.2 Comparison to results from papers

7.3 Result analysis

Chapter 8

Further work

Details all the results of your study here (exploits graphics for results visualisation). This chapter should also contain a full discussion, interpretation and evaluation of the results.

Chapter 9

Challenges

Summary

Details all the results of your study here (exploits graphics for results visualisation). This chapter should also contain a full discussion, interpretation and evaluation of the results.

9.1 Lack of datasets

9.2 Lack of universal comparison metric?

9.3 Academic papers algorithm description

Chapter 10

Project management

Summary

Details all the results of your study here (exploits graphics for results visualisation). This chapter should also contain a full discussion, interpretation and evaluation of the results.

10.1 Using GitLab

10.2 Canvas logs

10.3 other? Gantt chart?

Chapter 11

Conclusion

Conclusions should summarize the problem, the solution and its main innovative features, outlining future work on the topic or application scenarios of the proposed solution.

References

- [1] E. Weinstein and P. Moreno, “Music identification with weighted finite-state transducers,” in *2007 IEEE International Conference on Acoustics, Speech and Signal Processing-ICASSP’07*, vol. 2, pp. II-689, IEEE, 2007. 1
- [2] A. Wang *et al.*, “An industrial strength audio search algorithm,” in *Ismir*, vol. 2003, pp. 7–13, Washington, DC, 2003. 1
- [3] J. Haitsma, T. Kalker, and J. Oostveen, “Robust audio hashing for content identification,” in *International Workshop on Content-Based Multimedia Indexing*, vol. 4, pp. 117–124, Citeseer, 2001. 1
- [4] D. P. Ellis and G. E. Poliner, “Identifying cover songs’ with chroma features and dynamic programming beat tracking,” in *2007 IEEE International Conference on Acoustics, Speech and Signal Processing-ICASSP’07*, vol. 4, pp. IV-1429, IEEE, 2007. 2
- [5] T. Bertin-Mahieux and D. P. Ellis, “Large-scale cover song recognition using hashed chroma landmarks,” in *2011 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, pp. 117–120, IEEE, 2011. 2
- [6] D. P. Ellis and B.-M. Thierry, “Large-scale cover song recognition using the 2d fourier transform magnitude,” 2012. 2

REFERENCES

- [7] Spotify, “Music for everyone - spotify,” 2019. [Online; accessed 29-March-2019]. 2
- [8] Apple, “Music - apple (uk),” 2019. [Online; accessed 29-March-2019]. 2
- [9] SoundCloud, “Soundcloud – listen to free music and podcasts on soundcloud,” 2019. [Online; accessed 29-March-2019]. 2
- [10] Acoustic Glossary, “Sound power - acoustic glossary - article,” 2019. [Online; accessed 29-March-2019]. 4